

## CONTENT OF YOUR BASH .SH SCRIPT

3 parts, from the top of your .sh file:

1. '#!/bin/bash'
2. The SBATCH directives to Slurm
3. The job steps (example provided below: the job is to run an R script)

### 2. SBATCH directives to Slurm

```
#SBATCH --account=<myAccountName>    # Account to be charged for resources used
#SBATCH --job-name=<myJobName>        # Job name
#SBATCH --nodes=1                     # Nb of nodes to assign to this job (= on which to
distribute the Tasks)
#SBATCH --mem=<MB>                    # Memory required per node [in MB]
#SBATCH --time=10:00:00               # The time the job will take to run
(d-hh:mm:ss).
#SBATCH --mail-type=ALL                # Mail events (NONE, BEGIN, END,
FAIL, ALL)
#SBATCH --mail-user=<...@...>         # Email address to send job-related status
#SBATCH --output=</path/filename>.out # Set file in which to store output (default =
creates `slurm-#####.out` in working directory)
#SBATCH --error=<filename>            # Set file in which to store job error
messages
```

Optional:

```
#SBATCH -n x    Number of CPU cores you are requesting x = 5,10,20,etc
#SBATCH --ntasks=1                # max nb of tasks executed in // (across
nodes)
#SBATCH --ntasks-per-node         # nb simultaneous processes to run on each node
#SBATCH --cpus-per-task=1         # nb of CPU cores required per task
#SBATCH --mem-per-cpu=<MB>        # Memory required per allocated CPU core
#SBATCH --output=<filename>       # File in which to store job output
#SBATCH --array=<indexes>        # array range (ex: '--array=21-40%10')
```

Either --ntasks, or --ntasks-per-node & --nodes

Equivalent syntaxes:

- The account name for the job
  - #SBATCH --account=<myAccountName>
  - #SBATCH -A <myAccountName>
- The job name

- #SBATCH --job-name=<myJobName>
- #SBATCH -J <myJobName>
- set output & error files (default: creates `slurm-###.out` in wd)
- #SBATCH --output=<filename>
- #SBATCH --error=<filename>
- #SBATCH -o array\_%A.out
- #SBATCH -e array\_%A.error
- Total run time limit (d-hh:mm:ss)
- #SBATCH --time=10:00:00
- #SBATCH -t 10:00:00
- total number of tasks
- #SBATCH --ntasks=1
- #SBATCH -n 1
- nb of nodes
- #SBATCH --nodes=1
- #SBATCH -N 1

### 3. The job steps. Ex: R CMD BATCH

In this example, the computing 'job' is to run an R script.  
We will use the R tool 'BATCH', which is designed to run R in batch mode.

The general form will be:

module load R

R CMD BATCH <options> <infile> <outfile> <args>

#### <options>

default options are `--restore --save`

--save : saves all the objects in the workspace (i.e. all the objects created while the script was executed) into a hidden file `.RData`

--no-save : doesn't save data sets at the end of the R session

--restore : loads the content of the file `.RData` in the directory where R was started

--no-environ : doesn't read any user file to set environment variables

--vanilla : combines `--no-save`, `--no-environ`, `--no-init-file`, `--no-restore`

--silent : don't print out the initial copyright and welcome R messages.

#### <infile>

e.g. your. R script: `myscript1.R`

#### <outfile>

If no output file name `outfile` is provided, the name of `infile` is taken as default, appending the extension `.Rout` to it. I.e. here: `myscript1.Rout`. If you open this file, you will see the welcome message that appears everytime you open a new session in R, followed by the R commands (i.e. R code) that were executed, and finally an additional command at the end of the file with information about the execution time.